

AI-ENABLED SITE INSPECTION-TO-INVOICE RECONCILIATION

A Governance Framework for Verified, Auditable Progress Billing in EPC & Infrastructure Enterprises

AI Case Study — Smart Village — MBA — AI Track - 2026

Supervised by: Dr. Ahmed El-Seddawy

Ahmed Hassan Asfahany

Class 4A2

Abstract

Monthly progress invoicing on EPC and infrastructure contracts still depends, in most enterprises, on manually cross-checking site inspection findings against contractor claims — a process that is slow, sampled rather than exhaustive, and disconnected from the financial systems that ultimately pay the bill. This report presents a governance framework for AI-enabled inspection-to-invoice reconciliation, in which computer vision, natural-language processing, and machine-learning matching independently verify physical progress before an invoice is posted to the ERP, with robotic process automation (RPA) carrying verified data and routing decisions between systems. The framework is organized as a layered architecture — data capture, AI/analytics, integration/RPA, ERP, and monitoring — designed to keep judgment (what was actually built) decoupled from execution (how the record moves), so the model can scale across sites without being rebuilt each time. Applied to a representative EPC portfolio, the framework raises line-item review coverage from an estimated 15% to 100%, compresses reconciliation time from roughly 20.6 days to same-day, and is projected to catch approximately \$14.1M of a \$15.7M annual overbilling exposure before payment is made. The report also documents the framework's implementation limitations — legacy ERP connectivity, master-data quality, and change management — and outlines a path toward multi-site scaling and continuous model improvement.

Keywords: *AI-enabled reconciliation; ERP integration; Robotic Process Automation; Computer vision; Inspection-to-invoice; Responsible AI governance; EPC/infrastructure compliance*

Key Terms and Definitions

AI-Enabled Reconciliation — Using computer vision, NLP, and ML matching to independently verify billed quantities against physical progress before ERP posting.

BOQ (Bill of Quantities) — The contract schedule of unit rates and quantities against which site progress and invoices are measured.

RPA (Robotic Process Automation) — Software bots that move verified data and routing decisions between the AI layer and the ERP without manual re-entry.

Confidence Threshold — The model-certainty cut-off (this framework: $\geq 97\%$) above which a line item auto-approves rather than routing to human review.

Discrepancy / Overbilling Exposure — The dollar gap between what was invoiced and what AI verification confirms was actually completed.

Basis Points (bps) — 1/100th of one percent; used here to compare exposure across sites regardless of contract size (100 bps = 1%).

Class Imbalance — A training-data skew — e.g., far more “Pass” than “Fail” examples — that can cause a model to under-detect the rare class.

Discrete-Event Simulation — A modelling technique (used here in AnyLogic) that advances the model state at each discrete event — drone record, inspection, match — rather than continuously, to test process timing and bottlenecks.

I. Introduction

1.1 Background: The New EPC Enterprise Landscape Driven by AI

Engineering, Procurement, and Construction (EPC) and infrastructure contractors are under growing pressure to certify work faster without loosening financial controls. Quality-control (QC) teams capture physical progress through mobile inspection apps **ACC** (Autodesk Construction Cloud), photographs, and — increasingly — **drone** flights over active sites; Finance and Project Systems teams separately certify that progress against contract bills of quantities (BOQ) before releasing payment. In parallel, enterprise resource planning (ERP) platforms such as SAP and Oracle have absorbed more of the finance, procurement, and materials-management workflow, and robotic process automation (RPA) has become the standard way to move data between systems without manual re-keying. AI is now positioned to sit between these layers: not replacing inspectors or finance staff, but independently verifying that **what is billed matches what is actually on the ground**.

1.2 Problem: The Inspection-to-Invoice Gap

Despite this technology stack, a persistent, unquantified gap remains between site-quality inspection outcomes and what is submitted and paid through the monthly progress invoice. QC checklists, contractor claims, and ERP invoice records live in disconnected systems and are reconciled — if at all — through manual, sample-based review that can only cover a fraction of line items each cycle. The result is an operational gap (checkpoints do not talk to each other), a financial gap (unverified or over-stated quantities pass through to payment), and a strategic gap (disputes erode client trust and repeated audit findings damage credibility in future tenders).

1.3 Motivation: The Imperative for Verified, Auditable Billing

Closing this gap is not merely an efficiency exercise; it is a trust and sustainability imperative. Every line item that is billed but not independently verified is a latent financial exposure, a potential dispute, and — over enough cycles — a credibility risk with institutional and government clients who increasingly expect evidence-based certification. Conversely, a demonstrable, auditable verification process is a competitive differentiator: it shortens payment cycles, reduces dispute volume, and gives finance and audit teams a defensible record of why every dollar was approved. A systematic, AI-enabled approach to reconciliation is therefore essential to the long-term health of the contractor–client relationship.

1.4 Contributions

This report makes three contributions:

- A layered governance framework that embeds independent AI verification — computer vision, NLP/OCR, and ML matching — into the existing ERP-centred invoicing workflow, without requiring a rip-and-replace of financial systems.
- A practical toolkit for implementation: RPA bots for data movement, confidence-scored auto-approval thresholds, and dashboards that give QC, Finance, and Internal Audit a single verified view of project status.
- Evidence, drawn from a representative EPC portfolio, that standards-based, evidence-driven verification is not a drag on throughput but an enabler of faster, more defensible billing — full line-item coverage was reached without adding review headcount.

II. Background and Related Work

A. AI-Enabled Reconciliation: Definitions and Principles

AI-enabled inspection-to-invoice reconciliation refers to the use of computer vision, natural-language processing, and machine-learning matching to independently confirm that billed quantities correspond to verified physical progress, before that claim reaches the ERP's approval workflow. . **Four principles anchor the approach: 1- completeness (every line item is checked, not a sample); 2- traceability (every approval or exception carries the evidence and confidence score behind it); 3- timeliness (verification happens ahead of payment, not after a dispute); and 4- proportionate automation (only high-confidence matches auto-approve; everything else routes to a human reviewer).**

B. Enterprise Systems Landscape

Reconciliation does not happen in a vacuum — it sits at the intersection of three system layers already present in most EPC enterprises. Mobile QC/inspection apps (ACC) capture checklist results, photographs, GPS-tagged timestamps, and drone or BIM imagery at the point of inspection. ERP platforms (SAP PS/FI or Oracle Financials and Project Systems) hold the contract BOQ, budgets, and the accounts-payable register that ultimately reflects what was paid. RPA and middleware bridge the two, pulling ERP data on a schedule and posting verified or flagged results back without manual export/import. The framework in this report does not replace any of these systems; it adds a verification layer that reads from all three and writes a confidence-scored decision back to the ERP.

C. Governance and Standards Frameworks Relevant to Verified Billing

Several established governance and standards frameworks inform how this kind of AI verification layer should be built and audited, even though none was written specifically for construction billing:

- IEEE 7000-2021 — a model process for embedding ethical and human-values considerations into system design, applied here to how confidence thresholds and exception routing are set.
- ISO/IEC 42001:2023 — the AI management system standard, used as a reference for how model changes, monitoring, and retraining are governed over time.
- ISO 19011:2018 — guidelines for auditing management systems, informing how the framework's decision logs are structured to support internal and external audit.
- The COSO Internal Control – Integrated Framework — a widely used reference for financial control design, used to map the AI verification layer onto existing control objectives (authorization, accuracy, completeness) rather than treating it as a separate, parallel process.

D. Related Work

Table 1 situates the proposed framework against the reconciliation approaches already in common use across EPC and infrastructure portfolios.

Approach	Key Focus	Techniques / Methods	Limitations
Manual, sample-based QC audit	Spot-checking a fraction of invoices against inspection notes	Paper checklists, manual cross-referencing, monthly close review	Low coverage (~15%); slow (weeks); reconciliation is retrospective

Approach	Key Focus	Techniques / Methods	Limitations
Standalone RPA on the ERP	Automating data entry and invoice routing	Screen-scraping bots, scheduled ERP data pulls	Moves data faster but cannot judge whether the underlying claim is true
OCR / document automation only	Digitizing inspection reports and invoice PDFs	Optical character recognition, template-based field extraction	Extracts text but performs no independent verification against site evidence
Rules-based exception engines	Flagging invoices that breach fixed thresholds	Hard-coded tolerance bands, static approval rules	Brittle to new fraud patterns; generates high false-positive rates
Proposed AI-enabled reconciliation framework	Independently verifying physical progress against billed quantities, continuously	Computer vision, NLP/OCR, ML matching, drone imagery, RPA, confidence-scored auto-approval	Requires clean master data and change management; higher upfront build cost

Table 1: Related approaches to inspection-invoice reconciliation in EPC/infrastructure enterprises.


Across these approaches, a consistent pattern emerges: automation efforts to date have focused on moving data faster (RPA, OCR) or on catching known, fixed-threshold anomalies (rules engines), but none independently re-verifies the underlying physical claim at scale. This is the specific gap the proposed framework is designed to close.

III. Methodology: Proposed AI-Enabled Reconciliation Framework

The proposed methodology introduces a layered architecture in which each layer performs a distinct, complementary role — capturing site data, independently verifying it, moving the verified result, posting it to the system of record, and continuously monitoring the whole chain. This separation is deliberate: it keeps the judgment of “what was actually built” decoupled from the execution of “how the record moves,” so the framework can be extended to new sites or ERP instances without rebuilding the verification model each time.

7.2 — ARCHITECTURE

How AI Integrates with ERP Systems



Dashboards for Finance, PM, and QC teams draw from the ERP layer once data is posted — a single source of truth instead of three separate spreadsheets.

Figure 1: Layered architecture connecting site data capture, AI verification, RPA integration, and the ERP system of record.

At the base, the Data Capture Layer collects mobile QC-app checklists and contractor invoice PDFs as they are generated on site, including drone and BIM imagery where flights are scheduled. Above it, the AI/Analytics Layer applies computer vision to independently assess physical percentage-complete from photos, **drone** footage, and **BIM** overlays; NLP/OCR extracts structured quantities, cost codes, and budget figures from scanned reports and invoices; and machine-learning matching compares the two, flagging discrepancies and scoring its own confidence in each match. The Integration/RPA Layer then carries verified data and routing decisions to and from the ERP via API or middleware, with no manual export/import step. Finally, the ERP System Layer — Finance's accounts-payable register, Project Systems' BOQ and progress data, and Materials Management — becomes the shared source of truth that QC, Finance, and PM dashboards all read from, rather than three separate spreadsheets.

Because inspection formats, ERP fields, and dispute patterns evolve over the life of a portfolio, the framework is delivered and maintained through an iterative rather than a one-time build, illustrated in Figure 2.

SDLC Methodology

Agile / iterative SDLC recommended over Waterfall — inspection formats, ERP fields, and dispute patterns evolve, and the ML model needs continuous retraining rather than a one-time build.




Figure 2: Iterative delivery lifecycle for the AI verification model, from requirements through continuous maintenance.

Requirements are gathered jointly with QC and Finance stakeholders to define matching rules and exception thresholds; design maps the data architecture and model choice (computer vision, NLP, or ML) against the existing ERP; development proceeds in short sprints that build the extraction pipeline, train the matching model, and configure RPA bots; testing validates the model against historical invoices with known outcomes before go-live; deployment follows a phased rollout, piloting on one site before scaling to concurrent projects; and maintenance keeps the model current through continuous retraining on newly confirmed exceptions and a quarterly rule review — the same cadence that keeps the framework auditable as conditions change.

IV. Data Analytics: Where the Exposure Actually Lives

This section presents per-line diagnostic analysis of the underlying dataset (768 line items), starting with the most important finding: where the dollar exposure actually hides.




Figure 3: Share of total overbilling exposure by QC inspection outcome.




Figure 4: Invoiced vs. QC-verified amount per line item (linear scale, \$000s).




Figure 5: Invoiced vs. QC-verified amount per line item (log scale, \$000s).

The site-level risk concentration chart below surfaces the clearest actionable finding.




Figure 6: Overbilling exposure by site (bps of invoiced value) — SVC-03 is an outlier.

1. The biggest finding: exposure hides inside "Pass," not "Fail"

The chart above breaks the \$15.7M in overbilling exposure down by inspection outcome. Intuitively you'd expect it to concentrate in "Fail" or "Partial" items. It doesn't:

QC "Pass" lines carry 60.8% of the total dollar exposure (\$9.5M) even though they passed inspection.

"Partial" lines carry 38.9% (\$6.1M).

"Fail" lines carry only 0.2%.

This is statistically real, not noise: mean absolute discrepancy is 2.93% on "Pass" lines vs. 5.51% on "Partial/Fail" lines ($t = -11.6$, $p < 0.0001$) — so QC status is informative. But because 626 of 768 lines (81.5%) pass inspection, even a small average discrepancy on that huge base outweighs the larger average discrepancy on the small "Partial/Fail" base. A QC pass tells you the work is compliant — it doesn't tell you the billed quantity is correct. That's the gap a binary human sign-off structurally can't close, and it's the clearest evidence in this dataset for why line-item AI verification (not pass/fail sampling) is the thing worth automating.

2. Aggregate correlation is misleadingly reassuring

Invoice Amount and Verified Amount correlate at $r = 0.999$ in aggregate — if you only looked at totals, everything would look fine. The scatter chart shows why that's not good enough: at the line-item level, individual points still drift meaningfully off the diagonal, and it's specifically the off-

diagonal lines that carry the dispute risk. Portfolio-level correlation is the wrong altitude to catch this; only per-line comparison does.

3. The 20.6-day manual lag buys you nothing

Inspection Lag Days correlates with Discrepancy Pct at essentially zero ($r = -0.028$), and mean lag is almost identical for disputed vs. non-disputed lines (21.1 vs. 20.5 days). The delay isn't extra scrutiny — it's pure queue time. AI removing it isn't a speed/accuracy trade-off; there's no accuracy being traded away.

4. Root cause tells you which AI technique matters where

Among the 178 disputed lines, by dollar value:

Root cause	Lines	\$ value
Overstated Progress	79	\$5.58M
Change Order Not Reflected	18	\$1.89M
Quality Non-Compliance	36	\$1.74M
Measurement Error	45	\$1.58M

Table 2: Root cause of disputes, by line count and dollar value.

Overstated progress dominates — that's the computer-vision/drone-verification use case. Change-order mismatches are an NLP/document-matching problem against contract records. This is a real justification for the multi-technique (CV+NLP+ML) design rather than a single model.

5. One site is doing most of the damage

The site chart shows SVC-03 at 702 bps of exposure and a 53.1% dispute rate — 2–3x every other site (which cluster around 220–380 bps). This is the single most actionable finding: it says the framework shouldn't roll out uniformly, it should prioritize SVC-03 for immediate full-coverage review and a root-cause audit.

6. The cash-flow evidence for "prevent before payment"

\$23.95M is currently tied up across 94 unresolved "Under Review" lines — that's frozen cash under the manual process, today.

\$9.2M has already been clawed back after payment (84 lines) — meaning money physically left the door and had to be recovered. That's the exact failure mode pre-payment AI verification eliminates.

7. What this means for the threshold choice

A full threshold sweep confirms the trade-off directly: a 2% threshold flags 65.2% of lines (\$20.48M captured), while 5% (the framework's recommendation) flags 23.2% of lines but still captures \$10.79M — a better balance of catching value without burying reviewers (see Figure 14 in Section VIII). Combined with the confidence-band routing (50.1% auto-approve at $\geq 97\%$, 47.7% light review, 2.2% full review, 0% escalation), the practical AI benefit is precise: it doesn't need to review everything by hand — it needs to correctly triage everything, then only put human eyes on the ~50% that isn't already high-confidence.

4.5 Data Analytics Maturity: From Descriptive to Prescriptive

The dataset and workbook underpinning this report (Raw_Data, KPI_Dashboard, Predictive_Model, and Prescriptive_Scenarios sheets) span all four recognized levels of data-mining maturity. Classifying the analysis this way clarifies which questions each technique actually answers.

Descriptive (what happened) — KPI_Dashboard totals: \$630.0M invoiced across 768 line items, 8 sites, 12 months, 8 BOQ categories; 23.2% overall dispute rate.

Diagnostic (why it happened) — The seven findings in Section IV — e.g., 60.8% of exposure sits inside “Pass” lines, and SVC-03's dispute rate is 2–3x every other site — explain the root causes behind the descriptive totals.

Predictive (what happens next) — The workbook's Predictive_Model sheet scores each line item's overbilling risk from historical patterns, so a site/category's exposure can be estimated before the next invoice is even submitted.

Prescriptive (what to do about it) — The Prescriptive_Scenarios sheet and the threshold-sensitivity sweep in this section (Figure 14) turn that risk score into an action: which confidence threshold to set, and which line items to auto-approve versus route for review.

V. AnyLogic Process Simulation

The static analysis in Section IV establishes where exposure concentrates; it does not, by itself, show how the reconciliation process behaves over time — where items queue, how long rework loops take, or what share of “Pass” items still fail to match cleanly at the invoice stage. To test that dynamic behaviour, the framework's process flow was implemented as a discrete-event simulation in AnyLogic, using the Process Modeling library.

5.1 Model Architecture

The model mirrors the layered architecture from Section III at the level of individual line items. A Source block (“Drone”) generates inspection agents, which pass through a VisualInspection gate that routes each item to a Pass or Fail branch, matching the QC outcome probabilities observed in

the dataset. Pass items proceed through a processing delay (“Pass”) before reaching an Invoicesubmission gate, which represents the AI/ML matching step: items either match cleanly (“InvoiceMatch”, a terminal sink) or fail to match (“InvoiceNotmatch”), in which case they pass through a rework delay before being resubmitted (“Resubmit”). Fail items from the initial inspection route directly to a separate rework sink. This structure lets the simulation surface queueing and rework dynamics that a static spreadsheet cannot.




Figure 7: AnyLogic model canvas — Drone, VisualInspection, Pass/Fail branches, Invoicesubmission, and the InvoiceMatch/Resubmit/Rework sinks. Model compiled, 512 MB memory allocation, random seed 1.

5.2 Simulation Results

A representative run generated 173 inspection agents. 145 (83.8%) passed initial visual inspection and 28 (16.2%) failed — directionally consistent with the dataset's 81.5%/17.8% split once the residual “Partial” category is folded into a binary pass/fail gate. Of the 144 Pass items that completed processing and reached invoice submission, only 63 (43.8%) matched cleanly on the first pass; the remaining 81 (56.3%) required a rework cycle before 80 were successfully resubmitted. All 28 Fail items proceeded directly to rework.




Figure 8: Simulation run results — flow counts at each block (Drone → VisualInspection → Pass/Fail → Invoicesubmission → InvoiceMatch/InvoiceNotmatch → Resubmit/Rework).

This result reinforces, dynamically, the same finding Section IV established statically: passing visual inspection is not a reliable predictor of clean invoice matching. In this run, a majority (56.3%) of items that had already passed inspection still required a rework loop before their invoice matched — direct simulated evidence that “Pass” and “correct” are not the same event, and that a process relying on inspection status alone would let most of that rework happen after, rather than before, invoice submission.

VI. Case Studies and Applications

To demonstrate the practical relevance of the framework, this section traces how the layered architecture changes day-to-day work across the four business functions most directly affected by inspection-to-invoice reconciliation.




Figure 9: Business functions directly impacted by the AI-enabled reconciliation framework.

Quality Control

For QC teams, the framework does not change how inspections are performed — it changes what happens to the evidence afterward. Photographs, checklist results, and scheduled drone flights are no longer filed for manual, after-the-fact review; the computer-vision layer reads them as they arrive and cross-checks physical percentage-complete against the QC checklist and the BOQ. Drone imagery in particular gives the model a timestamped aerial record that is difficult to dispute, turning what was previously a supporting document into an active verification input.

Finance and Accounts Payable

Finance teams move from reconciling a sample of invoices against inspection notes once a month to reviewing only the line items the model actually flags. High-confidence matches (at or above the 97% threshold used in this framework) auto-approve; everything else is routed with the discrepancy and the evidence attached, so review time is spent on judgment calls rather than data assembly.

Procurement

Because verified quantities and confidence scores post directly to Materials Management and the contract BOQ, procurement planning is based on what has actually been installed or delivered rather than on self-reported contractor progress, reducing the risk of over-ordering against inflated completion claims.

Internal Audit

Every posting in the ERP now carries its AI confidence score and the inspection reference behind it. Instead of reconstructing a paper trail after the fact, Internal Audit can sample directly from a record that was built to be auditable from the moment the decision was made — a materially different starting point for both routine and regulatory audits.

VII. Experimental Evaluation and Comparative Analysis

The framework was evaluated against the baseline manual process on a representative EPC portfolio (total invoiced value: \$630.0M across the evaluated dataset), using four dimensions: line-item review coverage, reconciliation cycle time, overbilling exposure caught before payment, and audit-trail completeness.




Figure 10: Evaluation metrics comparison — baseline manual process vs. proposed AI-enabled framework.

Metric	Baseline (Manual)	Proposed Framework
Line-item review coverage	~15% (sampled)	100% (768/768 items)
Reconciliation cycle time	~20.6 days	Same day
Overbilling exposure caught pre-payment	~0% (found post-payment, if at all)	~90% (\$14.1M of \$15.7M identified)
Auto-approval rate $\geq$ 97% confidence	Not applicable	~50.1% of line items
Audit-trail completeness	Partial; reconstructed manually	Complete; logged at point of decision

Table 3: Comparative results, baseline vs. proposed framework.

Results show consistent improvement across every dimension measured. Line-item review coverage rises from an estimated 15% under manual sampling to 100% — all 768 line items in the evaluated dataset are checked, not a subset. Reconciliation cycle time compresses from roughly 20.6 days to same-day, because exceptions surface as data arrives rather than at month-end close. At a 5% dispute threshold, approximately 50.1% of line items qualify for automatic approval at 97% confidence or above, which is what allows full coverage without added review headcount. Most materially, of the \$15.7M in overbilling exposure identified across the dataset, the framework is projected to catch approximately \$14.1M — about 90% — before payment is made, compared with a baseline process that, by construction, only surfaces such exposure after the fact, if at all.



Figure 11: Financial impact and return-on-investment summary.

These figures translate directly into the financial case for the framework. Prevented overbilling protects revenue that would otherwise be disputed or written off; the shift from a ~20.6-day to a same-day reconciliation cycle shortens the cash-conversion cycle; and the combination of prevented overbilling and reduced manual-review labor is projected to cover the AI/RPA build cost within the first year of full rollout.

VIII. Discussion

Strengths

The framework's central strength is that it replaces sampling with completeness: every line item is checked using the same standard, rather than a fraction chosen by manual review capacity. This is paired with built-in traceability — every decision carries its confidence score and inspection reference — so audit readiness is a by-product of normal operation rather than a separate, after-the-fact exercise. The layered design also means QC, Finance, and Procurement continue to use their existing tools; the AI and RPA layers sit between systems rather than replacing them, which materially lowers the barrier to adoption.

Limitations

The framework also has real limitations. Legacy ERP systems often lack modern APIs, so older SAP/Oracle instances may need RPA-based “screen scraping” rather than clean REST integration. Master data quality is a persistent risk: inconsistent BOQ codes, units, or site IDs can break automated matching before it reaches the AI layer. Change management is non-trivial — finance and site teams accustomed to manual sign-off can be reluctant to trust an automated confidence score, particularly on higher-value line items. Finally, deep ERP customization or vendor-specific field mapping can be costly to build and maintain, especially for enterprises running multiple ERP versions across a portfolio.

These implementation limitations sit alongside three further risk categories — data privacy, algorithmic bias, and regulatory/adoption risk — summarized in Table 4 and visualized in Figures 12 and 13.

Risk Category	Description	Mitigation
Data Privacy & Security	Drone and site-photo imagery can capture identifiable workers and vehicles; QC records	Data minimization, role-based access, and encryption in transit/at rest, aligned to GDPR Art. 5/32 and ISO 27001 A.5.34 / A.8.24.

	link names to performance data (GDPR-scope personal data).	
Algorithmic Bias	Training data is 81.5% Pass / 17.8% Partial / 0.65% Fail — very few confirmed-fail examples; SVC-03 drives a disproportionate share of disputes.	Mandatory human review below the confidence threshold; quarterly bias audit comparing model flags against manual findings, by site.
Regulatory, Adoption & System Limits	Financial-audit and e-invoicing regulations; reviewer distrust of automated confidence scores; legacy ERP APIs and master-data quality.	Policy-as-code mapped to audit standards, phased rollout with human sign-off retained, RPA middleware to bridge legacy ERP APIs.

Table 4: Risk register — data privacy, algorithmic bias, and regulatory/adoption risk.




Figure 12: Risk heat map — likelihood vs. impact. Algorithmic Bias and Data Privacy & Security sit in the highest-priority zone.




Figure 13: Where the bias risk actually comes from — QC outcome class imbalance (left) and site-level exposure concentration (right).

Future Directions

Several extensions would strengthen the framework further. Multi-site rollout should allow the same verification model to extend across concurrent projects without a proportional increase in review headcount, provided the modular architecture described in Section III is maintained. Continuous learning — feeding confirmed exceptions back into training — should let matching accuracy compound over time rather than decay as invoice volume grows. Vendor-agnostic integration, keeping the RPA/middleware layer decoupled from any single ERP vendor, would reduce lock-in risk as the estate grows or changes. Looking further out, blockchain-enabled audit trails could give multiple stakeholders (contractor, client, and auditor) a shared, tamper-evident record without a centralized authority, and multimodal AI capable of reasoning jointly over drone video, BIM models, and structured ERP data would extend verification beyond what checklist-based computer vision can currently assess.

These directions translate into a phased roadmap (Table 5), and the choice of a 5% confidence-routing threshold is itself data-driven, as Figure 14 shows.

Phase	Focus	Key Actions
Now (0–3 mo)	Stabilize the outlier	Root-cause audit and full-coverage AI rollout at SVC-03; close the 94-line, \$23.95M “Under Review” backlog.
Next (3–9 mo)	Scale the core	Extend full-coverage verification to the remaining 7 sites; stand up the continuous-learning loop on confirmed exceptions.
Later (9–18 mo)	Decouple and extend	Move to a vendor-agnostic RPA/middleware layer; pilot multimodal verification joining drone video, BIM, and ERP data.
Vision (18 mo+)	Federate and share	Federated governance across the multi-site portfolio, with a blockchain-backed shared audit trail for clients and auditors.

Table 5: Phased implementation roadmap.




Figure 14: The data behind the recommendations — root cause of disputes by dollar value (left) and threshold sensitivity, coverage vs. dollars captured (right).

IX. Conclusion

This report has presented a layered AI-enabled reconciliation framework designed to close the persistent gap between site-quality inspection and monthly progress invoicing in EPC and infrastructure enterprises. By positioning computer vision, NLP, and ML matching as an independent verification layer ahead of ERP posting — connected through RPA rather than through a rebuild of existing financial systems — the framework raises review coverage from a manual sample to 100% of line items, compresses reconciliation time from weeks to same-day, and is projected to prevent approximately \$14.1M in annual overbilling exposure on the evaluated portfolio. The case studies across Quality Control, Finance, Procurement, and Internal Audit show that these gains do not come at the expense of existing workflows: each function keeps its tools and its judgment calls, while gaining a faster, better-evidenced basis for them. The framework's limitations — legacy ERP connectivity, master-data quality, and change management — are real but addressable, and the proposed path to multi-site scaling and continuous learning gives enterprises a concrete route from a single successful pilot to a portfolio-wide, auditable standard for verified billing.

References

1. IEEE Standards Association. (2021). IEEE 7000-2021: Model process for addressing ethical concerns during system design. Piscataway, NJ: IEEE.
2. IEEE Standards Association. (2020). IEEE 7010-2020: Recommended practice for assessing the impact of autonomous and intelligent systems on human well-being. Piscataway, NJ: IEEE.
3. International Organization for Standardization. (2023). ISO/IEC 42001:2023 — Information technology: Artificial intelligence management system.
4. International Organization for Standardization. (2018). ISO 19011:2018 — Guidelines for auditing management systems.
5. Committee of Sponsoring Organizations of the Treadway Commission (COSO). (2013). Internal Control – Integrated Framework.
6. Project Management Institute. (2021). A Guide to the Project Management Body of Knowledge (PMBOK Guide), 7th ed.
7. European Parliament and Council. (2016). Regulation (EU) 2016/679 (General Data Protection Regulation).
8. European Commission. (2021). Proposal for a regulation laying down harmonised rules on artificial intelligence (Artificial Intelligence Act).
9. The AnyLogic Company. (2024). AnyLogic simulation software (Version 8.9) [Computer software]. <https://www.anylogic.com>